




(a) trained on one distribution,  
judged against another



(b) locality kernel  
standard LIME



(c)  $\times$  class weights  
CIKM'23; 86% of the sample is one class



(d)  $\times$  density ratio  
pulls weight back towards the data

