




(a) trained on one distribution,
judged against another



(b) locality kernel
standard LIME



(c) $\times$ class weights
CIKM'23; 86% of the sample is one class



(d) $\times$ density ratio
pulls weight back towards the data

